


Deep Learning

What is Deep Learning

- Deep Learning is hierarchical feature learning
 - *The hierarchy of concepts allows the computer to learn complicated concepts by building them out of simpler ones. (Ian Goodfellow and Aaron Courville)*
 - *Deep learning allows computational models that are composed of multiple processing layers to learn representations of data with multiple levels of abstraction. (Yoshua Bengio and Geoffrey Hinton)*
- Deep Learning is large Neural nets
 - *When you hear the term deep learning, just think of a large deep neural net. (Jeff Dean)*

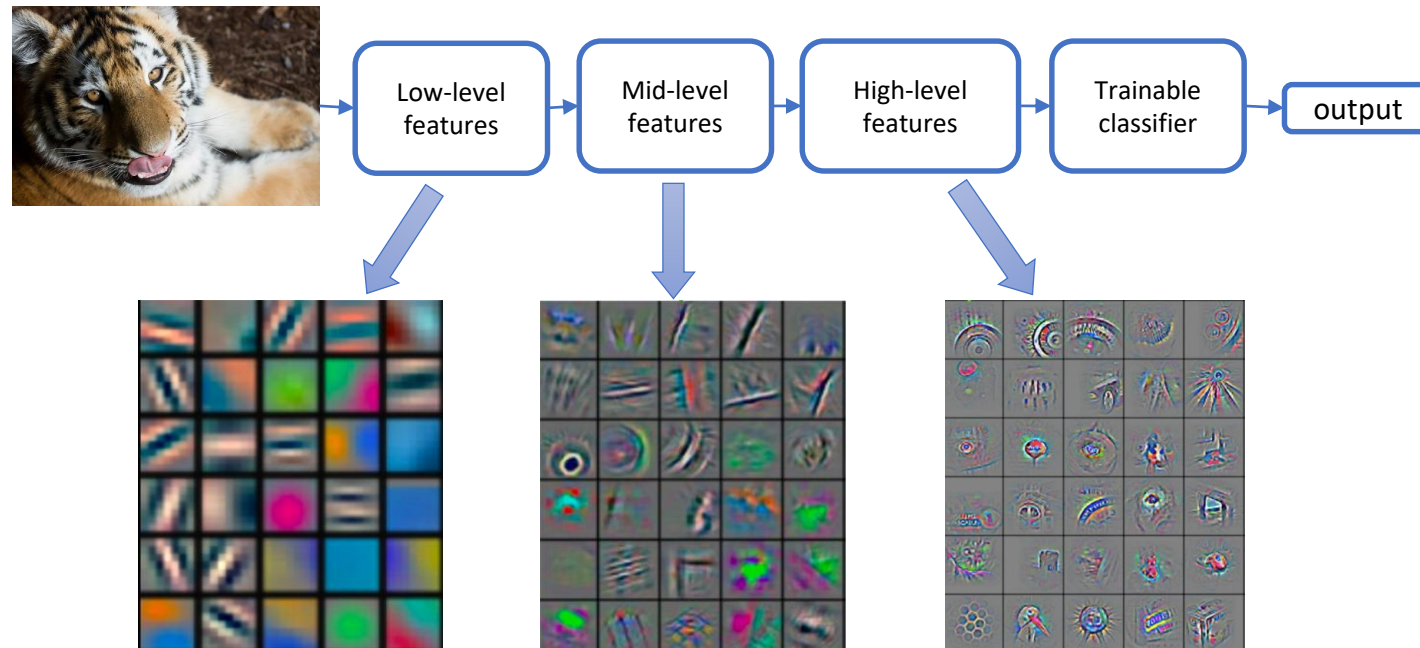
Universal approximation theorem



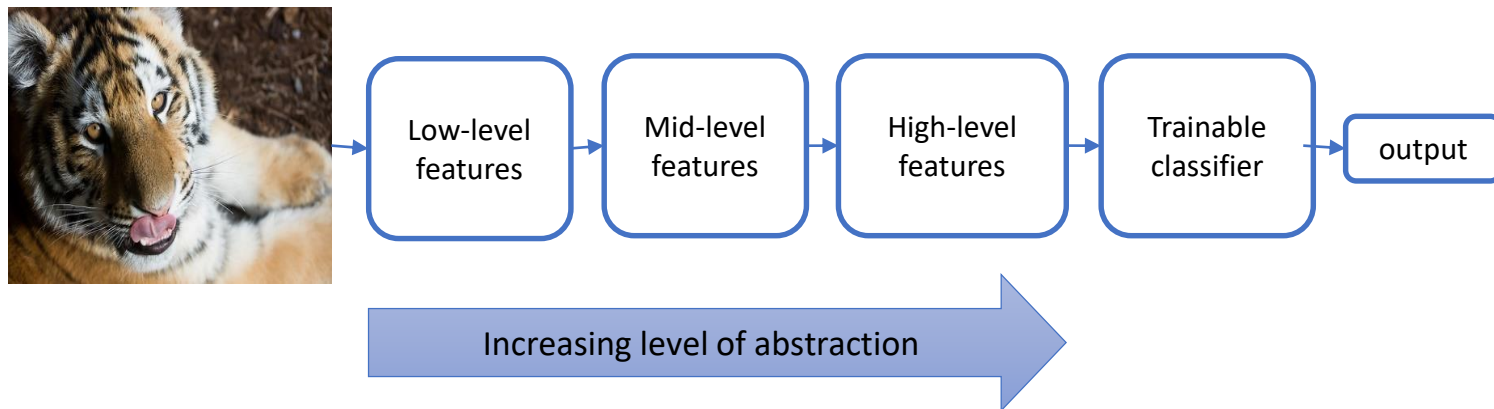
Theorem. A feedforward single hidden layer network with finite width can approximate continuous functions on compact subsets of \mathbb{R}^n under mild assumptions on the activation function.

Deep Nets as Generalizer

- Deep learning (a.k.a. representation learning) seeks to learn rich hierarchical representations (i.e. features) automatically through multiple stage of feature learning process.



Learning Hierarchical Representations



- Image recognition

- Pixel → edge → texon → motif → part → object


- Text

- Character → word → word group → clause → sentence → story

Why there is a sudden surge in interest

- Research on deep neural networks almost abandoned from 2000-2006
 - Overfitting, extremely slow training, local minima, gradient vanishing/exploding
- In 2006, Hinton, et. al. proposed RBMs to pretrain a deep neural network
- In 2009, Raina, et. al. proposed to use GPUs to train deep neural network
- In 2010, Dahl, et. al. trained a deep neural network using GPUs to beat the state-of-the-art in speech recognition
- In 2012, Krizhevsky, et. al. won the ImageNet challenge with NN
- In 2012, Mikolov, et. al. trained a recurrent neural network to achieve state-of-the-art in language modelling

Scale drives deep learning progress



Challenges of Deep Learning

- Deep Neural nets are data hungry
- Vanishing/exploding gradients
- Overfitting
- Hyperparameter Optimization
- Requires high-performance hardware

Vanishing/Exploding Gradient


$$\frac{\partial C}{\partial b_1} = \sigma'(z_1) \times w_2 \times \sigma'(z_2) \times w_3 \times \sigma'(z_3) \times w_4 \times \sigma'(z_4) \times \frac{\partial C}{\partial a_4}$$



Neural Network

- The sigmoid'(z1),sigmoid'(z2).. etc are less than ¼
- The weight matrices w1,w2,w3,w4 are initialized using gaussian method to have a mean of 0 and standard deviation of 1. Hence ||w(i)|| is less than 1
- If we initialize our weight matrices with very large values then gradient explodes and model becomes unusable

Restricted Boltzmann Machine




- Two-layered artificial neural network with generative capabilities
- Learns by minimizing reconstruction error

Pretraining DNN with RBMs

- The first layer RBM is trained using the data vectors v , resulting in parameters θ_1
- . In a subsequent layer k , the activation probabilities Q of the hidden units in the $k - 1$ 'th layer are calculated by propagating the data vectors v through the layers already learnt, and are used as the training vectors for the k 'th layer resulting in parameters θ_k
- This is greedy, layerwise and unsupervised pre-training.



Activation Functions

- Sigmoid –
 $(1 + e^{-x})^{-1}$



- Vanishing gradient
- Slow Convergence
- Saturate and kill gradients

- RELU – $\max(0, x)$



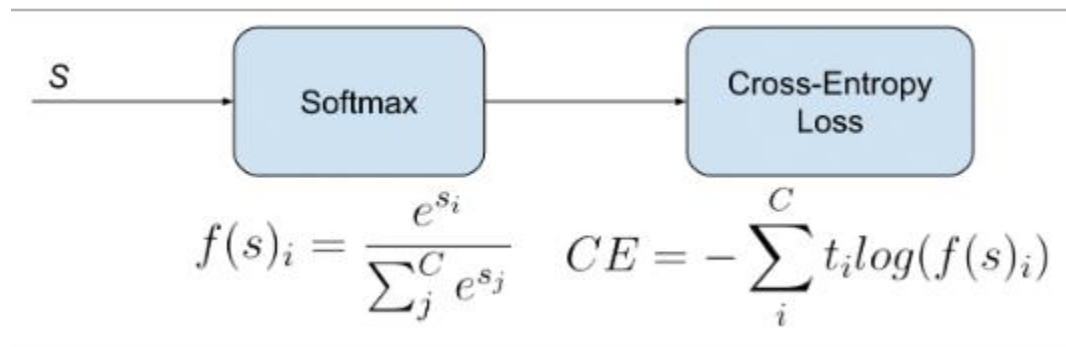
- Rectifies Vanishing gradient
- Leaky Relu resolves dead Neurons

- Softmax – $e^{z_j} / \sum_{k=1}^K e^{z_k}$




Loss and Cost Function

- MSE Loss = $\frac{1}{2} (\text{predicted value} - \text{actual value})^2$
- Cost function (J) = $1/m$ (Sum of Loss error for 'm' examples)
- Binary cross entropy = $-(Y \log(Y_{\text{predicted}}) + (1-Y) \log(1-Y_{\text{predicted}}))$
- Categorical cross entropy



Mini Batch Gradient Descent

- Divide n training dataset into m batches with batch size n/m . Typical batch size ranges from 64 – 512.
- Update weight after processing each batch
- If batch size = 1 \rightarrow Stochastic gradient descent
- Batch size = $n \rightarrow$ Batch gradient descent




- Batch gradient \rightarrow smooth but too long per iteration
- Stochastic gradient \rightarrow can avoid local optima but too noisy
- Batch gradient \rightarrow Fastest, avoid local optima, use vectorization less memory expensive

Gradient Descent with momentum

- Update weights with moving average of gradients
- V is momentum, gradient is acceleration and beta is friction

$$V_t = \beta V_{t-1} + (1 - \beta) \nabla_w L(W, X, y)$$
$$W = W - \alpha V_t$$



RMSProp and Adam

- Adapt learning rate by root mean square of moving average of gradient


$$E[g^2]_t = \beta E[g^2]_{t-1} + (1 - \beta) \left(\frac{\delta C}{\delta w} \right)^2$$

$$w_t = w_{t-1} - \frac{\eta}{\sqrt{E[g^2]_t}} \frac{\delta C}{\delta w}$$


- Adam is adaptive moment estimation and is combination of momentum gradient descent and RMSprop

Weight Initialization

- Random Initialization



$$\begin{aligned}\text{Var}(Z) &= 250 \\ \text{Std}(z) &= \sqrt{250} \\ &= 15.811\end{aligned}$$



- Xavier Initialization

- $\text{var}(\text{weights}) = \frac{1}{N}$
- $\text{weight} * \sqrt{\frac{1}{N}}$

- He Initialization



- $\text{weight} * \sqrt{\frac{2}{N}}$

Data Normalization



Why normalize inputs?

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Unnormalized:
 $w_1, x_1: 1 \dots 1000$
 $w_2, x_2: 0 \dots 1$



Normalized:



Andrew Ng

- Subtract mean and normalize variance

$$\bullet \quad x = x - \mu; \quad x = \frac{x}{\sigma^2}$$

Batch Normalization

- As the parameters of the preceding layers change, the distribution of inputs to the current layer such that the current layer needs to constantly readjust to new distributions
- Batch normalization accelerates training and allows using higher learning rate


For a layer with d dimensional input, each dimension is normalized

$$\hat{x}_i^{(k)} = \frac{x_i^{(k)} - \mu_B}{\sqrt{\sigma_B^{(k)^2} + \epsilon}}$$



To restore the representation, the mean and variance is readjusted, where gamma and beta are learnable parameters

$$y_i^{(k)} = \gamma^k \hat{x}_i^{(k)} + \beta^{(k)}$$

Covariate Shift



Cat $y = 1$ ✓ Non-Cat $y = 0$




$y = 1$ ✓ $y = 0$




"Covariate shift"

$\underline{x} \rightarrow y$

Why BatchNorm works




Learning rate decay



- Step decay with epochs
- Exponential decay with epochs

Overfitting in deep NN




- Complex Deep NN can memorize complex features
- Memorization is not learning!

Data Split

Small dataset



Big dataset




 Train set  Dev set  Test set



Data



Bad Split



Good Split

Detect Overfitting

	Low <i>Training</i> Error	High <i>Training</i> Error
Low <i>Testing</i> Error	The model is learning!	Probably some error in your code. Or you've created a <i>psychic</i> AI.
High <i>Testing</i> Error	OVERFITTING	The model is not learning.

Data Augmentation

- Gather more data
- Get diverse data
- Add noise



Original



1st Variation



2nd Variation



3rd Variation

Each iteration sees as **different variation** of the original sample.

L1/L2 Regularization

- $Cost = Loss + \frac{\lambda}{2m} ||w||^l \begin{cases} l = 2 \text{ for L2} \\ l = 1 \text{ for L1} \end{cases}$
- L2 regularization is also known as *weight decay* as it forces the weights to decay towards zero (but not exactly zero).
- Smaller weight leads to simpler model


L1:

$$w_{\text{new}} = \begin{cases} (w - H) - \lambda, & w > 0 \\ (w - H) + \lambda, & w < 0 \end{cases} \quad \begin{array}{l} \text{--- (1.1)} \\ \text{--- (1.2)} \end{array}$$

L2:

$$w_{\text{new}} = (w - H) - 2\lambda w \quad \text{--- (2)}$$

Dropout




- At every iteration, it randomly selects some nodes and removes them along with all of their incoming and outgoing connections
- Probability of choosing how many nodes to be dropped is another hyperparameter


Hyperparameter Tuning

- Learning rate
- Beta of RMSProp
- Hidden layers config
- Learning rate decay
- Mini batch size
- Regularization parameters

Grid Layout



Random Layout



Training a Deep NN

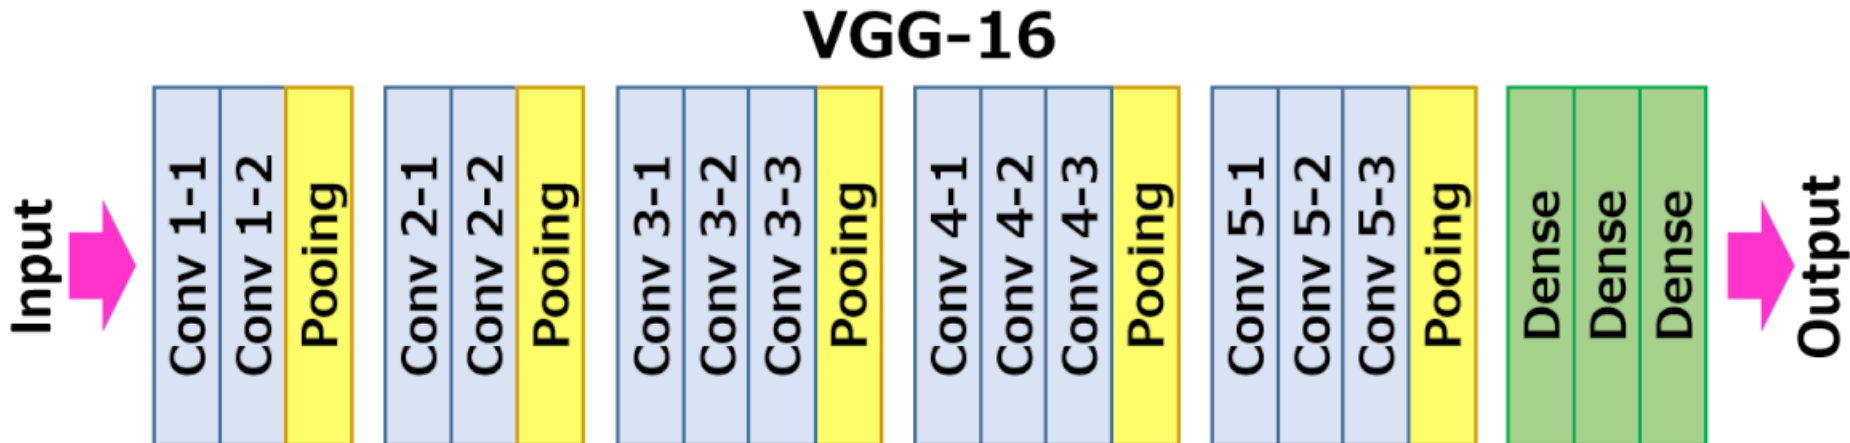
- Data preprocess
 - Data augmentation and normalization
 - Create training, validation and test set
- Design Network
 - Layers selection and configuration
 - Activation, Regularization and BatchNormalization
 - Weight Initialization
- Training model
 - Select optimizer like RMSProp, Adam etc.
 - Set hyperparameters
 - Train and record validation accuracy
- Fine tune and repeat
 - Tune hyperparameters and design and retrain to improve accuracy
 - Stop when required accuracy level reached

State of Art Architectures

- LeNet – for OCR

- Input -> Conv -> Relu -> Pool -> Conv -> Relu -> Pool -> FC -> Relu -> FC

- VGG Net



ResNet

- Motivation for skipping over layers is to avoid the problem of vanishing gradients, by reusing activations from a previous layer until the adjacent layer learns its weights
- Overcomes vanishing gradient for very deep nets
- Won Imagenet challenge on 2015




Figure 5. A deeper residual function \mathcal{F} for ImageNet. Left: a building block (on 56×56 feature maps) as in Fig. 3 for ResNet-34. Right: a “bottleneck” building block for ResNet-50/101/152.

Thank You